

Hierarchical Constraint-Aware Design-Space Exploration for Heterogeneous Distributed Cyber-Physical Systems

Nikita Tarasov^a

^a*Lviv Polytechnic National University, Lviv, Ukraine*

Abstract

Architecting a heterogeneous distributed cyber-physical system requires simultaneous decisions on the number and class of edge and gateway nodes, on where application stages are executed, on replication, and on communication links, under conflicting latency, energy, cost and network-volume objectives. The resulting design space grows combinatorially, and for distributed cyber-physical systems a single design point is expensive to evaluate, so exhaustive exploration is impractical while population-based search returns approximate fronts without guarantees. This paper presents HCA-DSE, a hierarchical constraint-aware design-space exploration method that discards structurally ill-formed, provably infeasible and provably dominated regions of the architecture space *before* any full design-point evaluation. Its central technical ingredient is a family of admissible objective bounds for *contention-dependent* objectives: the end-to-end latency of a distributed cyber-physical system depends on tier utilisation, which is not determined by any single design decision, so the assignment-monotonicity assumed by existing exact exploration methods does not hold. We prove that the resulting pruning is safe and that the exact Pareto front is preserved, and we verify both the bound admissibility and the front equality exhaustively on reference spaces. On four design-space scales, three workload classes and spaces of up to 8.3×10^7 candidate architectures, HCA-DSE performs 26–61 full evaluations instead of 10^3 to 10^7 , preserves the exhaustive Pareto front exactly (recall and precision of 1.0), and visits fewer than 2×10^{-4} of the candidates while never materialising the design space. Budget-matched NSGA-II requires 1000 evaluations to reach a median Pareto recall of 0.50–0.91. An evaluation-cost analysis identifies the break-even evaluator cost at 1–3 μ s, so the method pays off for any non-trivial evaluator; with a discrete-event evaluator the measured speed-up is up to 557 $\times$. Finally, the analytical model is validated against a containerised edge–gateway–cloud testbed with network shaping: over 36 architectures and 108 measured runs, the model orders architectures consistently with the measurements in 98% of the pairs it separates by more than 25%.

Keywords: design space exploration, cyber-physical systems, edge computing, multi-objective optimisation, Pareto optimality, architecture exploration, design automation

1. Introduction

Distributed cyber-physical systems (dCPS) are increasingly deployed on a heterogeneous compute continuum that spans sensing devices, edge nodes, gateways and cloud infrastructure. Designing such a system is not a single decision but a joint one: how many edge nodes and gateways

Email address: 7nikitan@gmail.com (Nikita Tarasov)

to deploy, which hardware class each of them should be, where each application stage is executed, how far the ingest stage is replicated, and which class of communication link connects the tiers. These decisions interact. Adding edge nodes reduces queueing delay but raises deployment cost and static energy; moving analytics to the cloud removes load from the edge tier but adds a wide-area transfer to every request; replication improves ingest throughput but multiplies both compute load and synchronisation traffic.

Because the decisions are discrete and interacting, the number of candidate architectures grows as the product of the decision domains and quickly reaches 10^5 to 10^7 even for a modest set of design variables. At the same time, evaluating a single candidate is expensive: a recent industrial dCPS exploration methodology reports that “evaluating each design point takes under a minute” with a discrete-event simulator, and its authors explicitly leave the integration of a search strategy out of scope [1]. A survey of the field concludes that “efficient and scalable DSE technology for dCPS is more or less non-existing and constitutes a largely uncharted research area” and identifies the cost of a design-point evaluation as a defining difficulty of the distributed setting [2]. The combination of a large space and an expensive evaluator is precisely the regime in which the number of *full evaluations*, rather than the raw size of the space, determines whether design-space exploration (DSE) is practical at all.

Two families of approaches dominate. Exhaustive or model-based enumeration evaluates every candidate and therefore yields the exact Pareto set, but its cost scales with the size of the space. Population-based metaheuristics such as NSGA-II [3] operate under a fixed evaluation budget and return an approximation with no guarantee on which trade-offs have been missed; as we show in Section 6, a budget-matched NSGA-II run needs 10^3 evaluations to recover half of the exact front on a space of 5×10^4 candidates. A third, smaller family performs *exact* multi-objective exploration by pruning partial designs, either symbolically with constraint solvers [4, 5] or through bound-based Pareto pruning in the broader optimisation literature [6, 7]. Methods in this family share one structural assumption: the objectives must be *assignment monotonous*, i.e. each additional decision can only add to the objective, so that a partial design already carries a valid lower bound on all its completions.

That assumption does not hold for the objective that matters most in a distributed CPS. End-to-end latency contains a waiting component that depends on the utilisation of a tier, and utilisation is a function of several decisions at once — the number of nodes in the tier, their class, the replication factor and the placement policy. No individual assignment increases latency monotonically, so the standard “cost so far plus optimistic remainder” bound is unavailable. This is the gap the present paper addresses.

Contributions.

1. A formal design-space model for heterogeneous edge–gateway–cloud CPS architectures (Section 3) with analytical models of latency including contention, energy, normalised deployment cost, network volume, and six classes of feasibility constraints.
2. **Admissible bounds for contention-dependent objectives** (Section 4). We decompose each objective into non-negative terms, bound every term independently over the still-reachable domains, and exploit a decision ordering under which the determinants of tier utilisation are fixed early, so that the waiting term becomes *exactly* computable on a partial design rather than being relaxed to zero. Lemmas 1 and 2 establish admissibility.
3. **HCA-DSE**, a hierarchical exploration method that combines structural well-formedness, feasibility pruning and bound-based dominance pruning with a final exact evaluation of the

survivors, together with proofs (Propositions 1–3, Theorem 1) that the exact Pareto front is preserved.

4. **An evaluation-cost analysis** that answers when such pruning is worth its own overhead, including a measured break-even evaluator cost and a crossover curve (Section 6).
5. **An empirical validation** on a containerised edge–gateway–cloud testbed with per-node link shaping: 36 architectures, 108 measured runs, reported both as prediction error and as the agreement between the architecture ordering produced by the model and the ordering observed in the measurements.
6. A reproducibility package in which every figure and every numeric table of this paper is generated from the raw result files by a single script, and in which the correctness claims are backed by exhaustive machine-checked tests.

The remainder of the paper is organised as follows. Section 2 positions the work. Section 3 defines the design space and the evaluation model. Section 4 presents HCA-DSE and its correctness. Section 5 describes the experimental methodology, Section 6 the results. Section 7 discusses when the method pays off and when it does not, Section 8 the threats to validity, and Section 9 concludes.

2. Related work

2.1. Exact multi-objective design-space exploration

The closest methodological relatives of this work perform exact multi-objective exploration by reasoning about partial designs. Neubauer et al. [4] encode system-level synthesis in Answer Set Programming modulo Theories and attach theory propagators that check both constraints and Pareto dominance on *incomplete* assignments, returning conflict clauses that remove infeasible and dominated regions early. Their formulation targets heterogeneous multiprocessors with routing, optimising latency, area and energy, and their pruning is sound for problems whose objectives are assignment monotonous. In the broader optimisation literature, bound-based Pareto pruning is a recognised family [6], and multi-objective A^* with admissible or consistent heuristics [8, 7] provides its theoretical foundation: a partial solution whose optimistic vector is dominated by a complete solution can be discarded without losing Pareto-optimal completions.

HCA-DSE belongs to this family, and we do not claim safe pruning as such as a contribution. The difference is the class of objectives. In a distributed CPS the dominant latency term is contention-dependent: it is a function of tier utilisation, which no single assignment determines. We therefore construct bounds that are admissible without assignment monotonicity, by combining a term-wise relaxation with an explicit decision ordering that makes the contention term exactly computable as soon as its determinants are fixed (Section 4). A second difference is operational: symbolic approaches evaluate design points inside the solver and are consequently evaluated with cheap objective functions, whereas we explicitly target the regime of an expensive external evaluator and quantify the evaluator cost at which pruning becomes worthwhile.

2.2. Design-space exploration for distributed cyber-physical systems

Herget et al. [2] survey DSE for dCPS and identify the two premises of this paper: design spaces that are far larger than in classical MPSoC exploration, and design-point evaluations that are considerably more expensive. CompDSE [1] follows the Y-chart philosophy [9] and derives workload, platform and mapping models from execution traces of an industrial lithography machine, evaluating design points with OMNeT++ [10]. It is an evaluator, not a search strategy: the authors

Table 1: Positioning with respect to the closest prior work.

Work	System class	Obj.	Exact front	Pruning nism	mecha-	Cont.- dep.	Exp. eval.	Real meas.
Neubauer et al. [4]	heterog. MP-SoC	3	yes	symbolic clauses on partial designs	conflict on partial	no	no	no
CompDSE [1]	industrial dCPS	perf.	—	none (no search)		yes	—	industrial traces
ContrArc [11]	CPS architectures	1	—	MILP + subgraph isomorphism		no	no	no
SPACE4AI-D [12]	edge-cloud continuum	1	no	heuristic search		partly	no	yes
NSGA-II family [3]	edge-cloud	2-4	no	none		partly	no	rarely
HCA-DSE (this work)	heterog. edge-cloud CPS	4	yes	hierarchical: struct. + feas. + dominance		yes	yes	yes

report a manual exploration of three variations and state that integrating search algorithms is out of scope. The two approaches are therefore complementary — CompDSE answers how a design point should be evaluated, HCA-DSE answers which design points need not be evaluated at all — and the per-design-point cost reported there falls in the regime where our crossover analysis predicts speed-ups of two orders of magnitude.

ContrArc [11] explores CPS architectures with a mixed-integer program combined with contract refinement checking, pruning through subgraph isomorphism. It minimises a single objective (cost) subject to requirements rather than constructing a Pareto front, and its pruning is constraint-based rather than bound-based.

2.3. Design-time placement in the edge-cloud continuum

SPACE4AI-D [12] is the closest work in terms of the problem setting: design-time selection of resources and component placement across edge, cloud and function-as-a-service layers. The problem is formulated as a mixed-integer non-linear program with a single objective (deployment cost) under response-time constraints, and solved with randomised greedy, local search, tabu search, simulated annealing and genetic algorithms; no exactness guarantee is provided. Its validation on a real deployment reports deviations of 12% in cost and 20% in execution time, which is a useful reference point for the accuracy reported in Section 6.

2.4. Metaheuristic exploration of edge-cloud deployments

A large body of work applies evolutionary and swarm metaheuristics to task allocation, offloading and service placement in edge-cloud systems, typically optimising combinations of latency, energy and cost. These methods are approximate by construction and their quality is a function of the evaluation budget. We use NSGA-II [3] as implemented in pymoo [13] as the budget-matched representative of this family, and report hypervolume [14], IGD^+ [15] and Pareto recall against the exhaustive ground truth. Our conclusion is not that such methods are inferior in general, but that when admissible bounds are available exactness is achievable at a small fraction of their evaluation budget.

Table 1 summarises the comparison.




Figure 1: Architecture template. Solid arrows show the default request path; dashed arrows show the alternative stage placements selected by the placement policy.

Table 2: Resource classes. Compute classes: capacity C [MIPS], memory M [MB], idle and peak power [W], cost [units]. Link classes: bandwidth B , propagation delay τ , transfer energy e , cost.

class	C / B	M / τ	P^{idle} / e	P^{peak}	cost
<i>compute node classes</i>					
edge-low	2000	2048	4	12	1
edge-med	6000	8192	8	30	2.4
edge-high	15000	16384	15	65	5.2
gw-low	4000	4096	6	20	1.8
gw-high	12000	16384	12	48	4
cloud-med	40000	131072	0	180	6
cloud-high	100000	262144	0	380	12
<i>link classes</i> (name, B [Mb/s], τ [ms], e [J/MB], cost)					
constrained	10	4	0.3	0.2	–
standard	100	2	0.12	0.6	–
fast	1000	1	0.05	1.4	–

3. System and design-space model

3.1. Architecture template

We consider the four-tier template of Fig. 1. A population of devices generates requests; each request traverses three application stages. The *pre-processing* stage τ_1 runs either on the edge tier or on the gateway tier, the *aggregation* stage τ_2 always runs on the gateway tier, and the *analytics* stage τ_3 runs either on the edge tier or in the cloud. Which of these variants is instantiated is a design decision (the placement policy).

A compute node class is a tuple $(C, M, P^{\text{idle}}, P^{\text{peak}}, c)$ of compute capacity [MIPS], memory [MB], idle and peak power [W] and normalised cost per node. A link class is a tuple $(B, \tau^{\text{prop}}, e, c)$ of bandwidth [Mb/s], propagation delay [s], transfer energy [J/MB] and cost per attached node. The gateway-to-cloud connection is fixed infrastructure ($B = 200$ Mb/s, $\tau^{\text{prop}} = 20$ ms) and is not a design variable. Table 2 lists the concrete classes used in the experiments.

3.2. Workload

A workload is $(n_d, \lambda, \{\tau_i\}_{i=1}^3, D, s^{\text{sync}})$ where n_d devices each emit requests at rate λ , giving a system arrival rate $\Lambda = n_d \lambda$; stage $\tau_i = (w_i, m_i, s_i)$ has computational work w_i [MI], resident memory m_i [MB] and inbound payload s_i [MB]; D is the end-to-end deadline and s^{sync} the per-request replica synchronisation payload. Three workload classes are used throughout (Table 3); they are qualitatively different rather than numerically rescaled versions of one another: telemetry is high-rate and small-payload, real-time control is deadline-dominated, and data-intensive sensing is payload- and compute-dominated.

Table 3: Workload classes. Stage values are given as $\tau_1/\tau_2/\tau_3$.

workload	n_d	λ	Λ	w_i [MI]	s_i [kB]	D [ms]	$C^{\max}$
telemetry	200	2	400	8/4/20	20/10/10	120	60
control	60	10	600	5/3/9	5/4/3	30	60
sensing	40	1.5	60	120/30/260	1200/350/300	900	90

Table 4: Design variables and domain sizes of the four design-space scales.

variable	meaning	S1	S2	S3	S4
n_e	number of edge nodes	3	6	24	80
$edge_cls$	edge node class	3	3	3	3
n_g	number of gateways	2	3	5	8
gw_cls	gateway class	2	2	2	2
$link_cls$	access link class	3	3	3	3
$cloud_cls$	cloud instance class	2	2	2	2
rep	replication factor of stage 1	2	3	4	5
$policy$	placement policy	3	3	3	3
<i>design-space size</i> $ \mathcal{X} $		1 296	5 832	51 840	345 600

3.3. Decision vector and design space

An architecture is described by the vector

$$x = (n_e, \kappa_e, n_g, \kappa_g, \ell, \kappa_c, r, \pi), \quad (1)$$

with n_e, n_g the number of edge nodes and gateways, $\kappa_e, \kappa_g, \kappa_c$ the edge, gateway and cloud classes, ℓ the access link class, r the replication factor of stage τ_1 and π the placement policy. The design space is the Cartesian product $\mathcal{X} = D_1 \times \dots \times D_k$ with $k = 8$, and $|\mathcal{X}| = \prod_i |D_i|$. Table 4 lists the domains of the four scales used in the experiments.

3.4. Evaluation model

Let $t(i)$ denote the tier executing stage i under policy π , N_t the number of nodes of tier t and C_t the per-node capacity of its class. The offered load and utilisation of a tier are

$$\Theta_t = \sum_{i: t(i)=t} \Lambda w_i \mu_i, \quad \rho_t = \frac{\Theta_t}{N_t C_t}, \quad \mu_i = \begin{cases} r & i = 1 \\ 1 & \text{otherwise,} \end{cases} \quad (2)$$

that is, replication multiplies the load of the ingest stage but not of the others.

Latency. Stage service is deterministic, so a tier behaves as an M/D/1 queue and the Pollaczek–Khinchine formula [16] gives the waiting time. With $\mathcal{H}(x)$ the set of communication hops implied by the policy and the replication factor,

$$L(x) = \underbrace{\sum_{i=1}^3 \frac{w_i}{C_{t(i)}}}_{\text{service}} + \underbrace{\sum_{i=1}^3 \frac{w_i}{C_{t(i)}} \cdot \frac{\rho_{t(i)}}{2(1 - \rho_{t(i)})}}_{\text{waiting}} + \underbrace{\sum_{(s,\ell) \in \mathcal{H}(x)} \left(\frac{8s}{B_\ell} + \tau_\ell^{\text{prop}} \right)}_{\text{communication}}. \quad (3)$$

The middle term is what makes the objective *contention-dependent*: it is a function of ρ_t , which by (2) depends jointly on n_t , the class of the tier, the replication factor and the policy. If $\rho_t \geq 1$ for any used tier the architecture is discarded as unstable.

Energy, cost and network volume. Over an accounting horizon T ,

$$E(x) = \sum_{t \in \{e, g\}} N_t \left(P_t^{\text{idle}} + \rho_t (P_t^{\text{peak}} - P_t^{\text{idle}}) \right) T + \rho_c P_c^{\text{peak}} T + \sum_{(s, \ell) \in \mathcal{H}(x)} s \Lambda e_\ell T, \quad (4)$$

$$C(x) = n_e (c_{\kappa_e} + c_\ell) + n_g (c_{\kappa_g} + c_\ell) + c_{\kappa_c}, \quad (5)$$

$$D(x) = \Lambda \sum_{(s, \ell) \in \mathcal{H}(x)} s. \quad (6)$$

The objective vector is $F(x) = (L(x), E(x), C(x), D(x))$, all components minimised.

Feasibility. An architecture is feasible when all of the following hold: $\rho_t \leq \rho^{\max}$ for every used tier; the resident memory of the stages mapped to a tier fits in the memory of its node class; the aggregate access-link traffic does not exceed the total bandwidth of the attached nodes; $L(x) \leq D$; $C(x) \leq C^{\max}$; and $E(x) \leq E^{\max}$. We write $\mathcal{F}$ for the set of feasible architectures.

4. HCA-DSE

4.1. Partial architectures

The exploration traverses a prefix tree over the decision vector (1) under a fixed variable order $v_1, \dots, v_k$.

Definition 1 (Partial architecture). A partial architecture of depth d is a prefix of assignments $p = (a_1, \dots, a_d)$, $a_i \in D_i$. Its completions are $\mathcal{C}(p) = \{x \in \mathcal{X} : x_i = a_i, i \leq d\}$, with $|\mathcal{C}(p)| = \prod_{i>d} |D_i|$.

Let $\sigma : \mathcal{X} \rightarrow \{0, 1\}$ be the well-formedness predicate of the template: the number of gateways does not exceed the number of edge nodes, the replication factor does not exceed the number of edge nodes, and replication is undefined for policies that do not place the ingest stage on the edge tier. The working space is $\mathcal{X}_\sigma = \{x : \sigma(x) = 1\}$ and all statements below refer to $\mathcal{F}_\sigma = \mathcal{F} \cap \mathcal{X}_\sigma$; the exhaustive baseline against which we compare uses the same predicate. The Pareto front is

$$\mathcal{P}^* = \{F(x) : x \in \mathcal{F}_\sigma, \nexists y \in \mathcal{F}_\sigma : F(y) \preceq F(x)\}, \quad (7)$$

where $u \preceq v$ denotes $u_i \leq v_i$ for all i and $u_j < v_j$ for at least one j .

4.2. Admissible bounds

Definition 2 (Admissibility). LB_i is admissible for objective f_i if $LB_i(p) \leq f_i(x)$ for every partial architecture p and every $x \in \mathcal{C}(p)$.

Exact exploration methods that prune partial solutions assume that the objectives are *assignment monotonous*, i.e. that every assignment can only increase the objective, so that the partial cost is itself a bound [4]. The waiting term of (3) violates this assumption: it depends on ρ_t , which by (2) is jointly determined by n_t , the class of the tier, r and π , and which *decreases* when nodes are added. We therefore construct the bounds in two steps.

Lemma 1 (Term-wise relaxation). Let $f_i(x) = \sum_{t \in T_i} \varphi_t(x)$ with $\varphi_t \geq 0$, and let

$$\underline{\varphi}_t(p) = \min_{x \in \mathcal{C}(p)} \varphi_t(x).$$

Then $LB_i(p) = \sum_{t \in T_i} \underline{\varphi}_t(p)$ is admissible.

Proof. For every $x \in \mathcal{C}(p)$ and every t , $\underline{\varphi}_t(p) \leq \varphi_t(x)$ by definition of the minimum. Summing over $t \in T_i$ yields $LB_i(p) \leq f_i(x)$. $\square$

The minima are taken *independently per term*, so the combination of values that realises them need not correspond to any single completion; the bound is therefore optimistic but valid. Every term of (3)–(6) is monotone in each of its free variables — service time decreases in C , transfer time decreases in B , static energy and cost increase in the node count, network volume increases in r — so each $\varphi_t(p)$ is attained at a domain endpoint and is obtained from a table of precomputed domain extrema in $O(1)$ time rather than by search.

Lemma 2 (Exactness on determined terms). *Let $\text{supp}(\varphi_t)$ be the set of variables φ_t depends on. If $\text{supp}(\varphi_t) \subseteq \{v_1, \dots, v_d\}$ then φ_t is constant on $\mathcal{C}(p)$, and substituting its exact value for $\underline{\varphi}_t(p)$ preserves admissibility while tightening the bound.*

Proof. Constancy is immediate, and a constant equals its own minimum over $\mathcal{C}(p)$, so Lemma 1 applies unchanged. $\square$

Remark 1 (Why the variable order matters). The order $v_1, \dots, v_k$ places the determinants of tier utilisation — node counts, node classes, replication factor and policy — before the remaining variables. As soon as they are assigned, the waiting term

$$W_t = \frac{w_i}{C_t} \cdot \frac{\rho_t}{2(1 - \rho_t)}$$

is constant on $\mathcal{C}(p)$ and enters the bound *exactly* by Lemma 2, instead of being relaxed to its trivial bound 0. This is what makes a contention-dependent objective usable for safe pruning without assignment monotonicity. If the determinants are fixed and $\rho_t \geq 1$, no completion has finite latency and $LB_L(p) = \infty$, which is admissible.

4.3. Pruning rules

Proposition 1 (Safe feasibility pruning). *Let LB_{g_j} be admissible for constraint function g_j . If $LB_{g_j}(p) > g_j^{\max}$ then $\mathcal{C}(p) \cap \mathcal{F} = \emptyset$.*

Proof. For any $x \in \mathcal{C}(p)$, $g_j(x) \geq LB_{g_j}(p) > g_j^{\max}$, hence $x \notin \mathcal{F}$. $\square$

Proposition 2 (Safe dominance pruning). *Let $y \in \mathcal{F}_\sigma$ be a fully evaluated feasible architecture with $F(y) \preceq LB(p) = (LB_1(p), \dots, LB_n(p))$. Then $F(x) \notin \mathcal{P}^*$ for every $x \in \mathcal{C}(p)$.*

Proof. Let $x \in \mathcal{C}(p)$. By admissibility $f_i(x) \geq LB_i(p) \geq F_i(y)$ for all i , and by assumption there is a j with $F_j(y) < LB_j(p) \leq f_j(x)$. Hence y strictly dominates x and $F(x)$ is not on the Pareto front. $\square$

Proposition 3 (Closure of the well-formedness predicate). *Each invariant of σ depends only on assigned variables and, under the chosen order, all variables of an invariant are assigned before it is first checked. Since assignments are never revised, $\neg\sigma(p)$ implies $\sigma(x) = 0$ for every $x \in \mathcal{C}(p)$, i.e. $\mathcal{C}(p) \cap \mathcal{X}_\sigma = \emptyset$.*

Algorithm 1 HCA-DSE

Require: design space $\mathcal{X}$ with order $v_1..v_k$, workload w , evaluator F

Ensure: Pareto front $\mathcal{P}$

```
1:  $A \leftarrow \emptyset$  ▷ archive of non-dominated evaluated vectors
2:  $S \leftarrow \emptyset$  ▷ evaluated feasible architectures
3: procedure EXPLORE( $p, d$ )
4:   if  $d = k$  then
5:      $(F(p), \text{feasible}) \leftarrow \text{EVALUATE}(p, w)$  ▷ the expensive step
6:     if  $\text{feasible}$  then  $S \leftarrow S \cup \{p\}$ ; ARCHIVEINSERT( $A, F(p)$ )
7:   end if
8:   return
9: end if
10: for all  $a \in D_{d+1}$  do
11:    $q \leftarrow p \cup \{v_{d+1} \mapsto a\}$ 
12:   if  $\neg \sigma(q)$  then continue ▷ Prop. 3
13: end if
14:    $LB \leftarrow \text{BOUNDS}(q)$  ▷ Lemmas 1, 2
15:   if  $\exists j : LB_{g_j}(q) > g_j^{\max}$  then continue ▷ Prop. 1
16: end if
17:   if  $\exists f \in A : f \preceq LB$  then continue ▷ Prop. 2
18: end if
19:   EXPLORE( $q, d + 1$ )
20: end for
21: end procedure
22: EXPLORE( $\emptyset, 0$ )
23: return non-dominated subset of  $\{F(x) : x \in S\}$ 
```

Remark 2. Structural invariants define what counts as an architecture rather than following from the constraints of the problem: the evaluator can formally compute objectives for, say, a replicated deployment under a policy that does not place the ingest stage on the edge tier, although such a configuration is meaningless. Structural pruning is therefore part of the definition of the search space, and the exhaustive baseline applies the same predicate. This distinction was found by the test described in Section 5, which initially — and incorrectly — verified structural pruning as a feasibility argument.

4.4. Algorithm

Algorithm 1 traverses the prefix tree depth-first. At every node the three rules are applied in increasing order of cost: structural validity, then feasibility bounds, then dominance against an archive of non-dominated evaluated points. A leaf that survives is evaluated in full; feasible leaves update the archive. Fig. 2 illustrates the two bound-based rules.

4.5. Correctness

Theorem 1 (Pareto-front preservation). *If all bounds are admissible, pruning is performed only by Propositions 1–3, and the archive contains only feasible fully evaluated architectures, then the set of non-dominated vectors returned by Algorithm 1 equals the true Pareto front: $\mathcal{P}^{\text{HCA}} = \mathcal{P}^*$.*




Figure 2: Pruning on the prefix tree. Sub-trees are discarded when the optimistic bound already violates a constraint (Proposition 1) or is dominated by an evaluated feasible design (Proposition 2); only surviving leaves reach the expensive evaluator.

Proof. ($\subseteq$) Every returned vector is $F(x)$ of an evaluated $x \in \mathcal{F}_\sigma$ and survives non-dominated filtering; a dominating architecture would itself have been evaluated or pruned only by Proposition 2, which is excluded by that proposition.

($\supseteq$) Let $u \in \mathcal{P}^*$ and let $x \in \mathcal{F}_\sigma$ with $F(x) = u$. The traversal is a prefix tree, so x is reached through the unique chain $p_0 \subset p_1 \subset \dots \subset p_k = x$. Suppose some p_d were pruned. By Proposition 3 this implies $\sigma(x) = 0$ and by Proposition 1 $x \notin \mathcal{F}$, both contradicting $x \in \mathcal{F}_\sigma$; by Proposition 2 it implies the existence of a feasible y strictly dominating x , contradicting $u \in \mathcal{P}^*$. Hence no p_d is pruned, x is evaluated, and u survives the final filtering. $\square$

Remark 3 (Duplicates). The theorem is stated in objective space. If two distinct architectures share an objective vector, dominance pruning may remove one of them; the front is unaffected but the duplicate solution is lost. All experiments therefore compare sets of objective vectors.

Remark 4 (Order independence). Proposition 2 holds for any archive of evaluated feasible architectures, so the traversal order affects pruning efficiency but never correctness.

4.6. Complexity

Let C_E be the cost of one full evaluation, C_P the cost of computing the bounds at one tree node, N the number of visited partial states, N_E the number of full evaluations and $|A|$ the archive size. Then

$$T_{\text{exh}} = O(|\mathcal{X}| C_E), \quad T_{\text{HCA}} = O(N(C_P + n|A|) + N_E C_E), \quad (8)$$

and pruning is profitable when

$$C_E > \frac{N(C_P + n|A|)}{|\mathcal{X}| - N_E}. \quad (9)$$

Section 6 measures the right-hand side and shows that it lies in the microsecond range, and that $N \ll |\mathcal{X}|$, so the Cartesian product is never materialised.

5. Experimental methodology

5.1. Research questions

RQ1 By how much does HCA-DSE reduce the number of *full design-point evaluations*, and how is that reduction distributed over the pruning layers?

RQ2 Is the exact Pareto front preserved with respect to exhaustive enumeration?

RQ3 At which evaluator cost does the pruning overhead pay for itself?

RQ4 How does HCA-DSE compare with budget-matched stochastic exploration?

RQ5 How do the reduction and the runtime scale with the size of the design space?

RQ6 Do the architecture rankings produced by the analytical model transfer to a real containerised edge–cloud deployment?

5.2. Design of the experiments

Four design-space scales (S1–S4, Table 4) are combined with the three workload classes of Table 3. Exhaustive ground truth is computed for S1–S3; S4 and the scalability study are explored by HCA-DSE only. Every configuration is deterministic except the stochastic baselines, which are run with 20 independent seeds.

Baselines. Exhaustive enumeration provides the ground-truth Pareto front. Random search and NSGA-II [3] (pymoo [13], population 40, SBX crossover, polynomial mutation, integer repair) are run under matched evaluation budgets of 100, 250, 500 and 1000 full evaluations, 20 seeds each. All methods use exactly the same evaluator and the same feasibility rules, so differences are attributable to the exploration strategy alone.

Metrics. Pareto recall and precision are computed on sets of objective vectors against the exhaustive front. Hypervolume [14] is normalised by the hypervolume of the exhaustive front (HV ratio); IGD^+ [15] uses the exhaustive front as reference. We additionally report the number of full evaluations N_E , the number of visited partial states N , and wall-clock time. Stochastic baselines are compared with the Mann–Whitney U test [17] and Cliff’s δ [18].

Ablation. Four variants isolate the contribution of each layer: structural only; structural + feasibility bounds; structural + dominance bounds; and the full method.

5.3. Two-level evaluation and the cost model

The analytical evaluator of Section 3 (level A) is used for pruning, ground truth and all exactness claims. A discrete-event simulator (level B) provides a higher-fidelity, and therefore more expensive, evaluator: multi-server queues with a finite number of servers per tier, explicit link contention, deterministic service times and Poisson arrivals. Its measured cost is 19.4 ms per design point against $\approx 8 \mu\text{s}$ for the analytical evaluator, which allows the crossover analysis of RQ3 to be anchored at two measured points rather than assumed.

5.4. Testbed

The physical validation runs the architecture template as containers (Fig. 3). Each edge node, gateway and cloud instance is a separate container limited to one CPU, so that a node behaves as the single server assumed by the model. Services are implemented in Go and share one binary: a request carries a plan listing the remaining stages, and each node consumes CPU work equal to w_i/C_t before forwarding the remainder together with the real payload to the next hop. Links are shaped per node with `tc netem` [19] using the delay and rate of the link class of the design under test; the gateway-to-cloud hop uses the WAN profile. Replication is emulated faithfully: r replicas execute stage τ_1 in parallel and carry the synchronisation payload, while the remainder of




Figure 3: Containerised testbed. Every tier node is a separate one-CPU container; access links carry the delay and bandwidth of the link class of the design, the cloud hop carries the WAN profile.

the chain is executed once. A load generator produces Poisson arrivals at rate Λ , 900 requests per run of which the first 150 are discarded as warm-up, and reports the latency percentiles.

Thirty-six architectures were selected across the three workloads — Pareto-optimal, dominated and near-boundary designs — restricted to those that fit on a single machine (≤ 4 edge nodes, ≤ 2 gateways, one cloud instance). Each was executed three times, giving 108 measured runs.

Because the analytical model is deliberately coarse, we report both the prediction error (MAPE, median absolute error) and, more importantly for exploration, the *ordering* agreement: Spearman’s ρ , Kendall’s τ , and the share of architecture pairs whose measured order agrees with the model order among the pairs that the model separates by more than 10% and 25%.

5.5. Reproducibility

The implementation is in Python 3 with NumPy, SciPy and pymoo; the testbed services are in Go and run under Docker Compose. All figures and all numeric tables of this paper are generated from the raw result files by a single script, so no number is transcribed by hand. Correctness is machine-checked at three levels: (i) bound admissibility is verified *exhaustively* on S1, i.e. for every prefix and every completion of that prefix across the three workloads; (ii) for every pruned prefix, all of its completions are enumerated and the claim of the corresponding proposition is verified directly, on S1 and S2; and (iii) the equality $\mathcal{P}^{\text{HCA}} = \mathcal{P}^*$ is re-checked on S1–S3 for every pruning variant as a regression test.

6. Results

6.1. RQ1: most architectures are discarded before evaluation

Table 5 and Fig. 4 report the exploration funnel. On S3 (51 840 candidates) HCA-DSE performs 41 to 54 full evaluations depending on the workload, a reduction of 99.90% with respect to the raw space; on S4 (345 600 candidates) it performs 52 to 60, a reduction of 99.98%. The composition of the funnel differs by workload: for data-intensive sensing the feasibility bounds remove the largest share, whereas for telemetry dominance pruning does most of the work, which reflects the different binding constraints of the two regimes.

6.2. RQ2: the exact front is preserved

For every scale for which exhaustive enumeration is tractable (S1–S3) and for every workload, the set of objective vectors returned by HCA-DSE is identical to the exhaustive Pareto front: recall and precision are both 1.0. The same holds for each ablated variant, confirming that no individual pruning layer is unsafe. Bound admissibility was verified exhaustively on S1 — all prefixes against all completions — without a single violation, and for every pruned prefix on S1 and S2 all completions were enumerated and found consistent with the claim of the corresponding proposition. Fig. 5 shows the resulting front for S3/telemetry in two projections.

Table 5: Exploration funnel. Reduction is with respect to the raw design space; recall is against exhaustive enumeration where the latter is tractable.

scale	workload	$ \mathcal{X} $	structural pruned	feasibility pruned	dominance pruned	full evals	feasible	$ \mathcal{P} $	reduction [%]	exhaustive evals	recall
S1	telemetry	1 296	138	233	893	32	25	17	97.53	1 080	1.0
S1	control	1 296	138	269	863	26	23	18	97.99	1 080	1.0
S1	sensing	1 296	77	613	590	16	11	11	98.77	1 080	1.0
S2	telemetry	5 832	1 422	1 127	3 235	48	40	27	99.18	3 888	1.0
S2	control	5 832	1 425	1 236	3 133	38	35	25	99.35	3 888	1.0
S2	sensing	5 832	1 075	2 323	2 408	26	19	16	99.55	3 888	1.0
S3	telemetry	51 840	8 374	27 573	15 839	54	46	31	99.90	34 560	1.0
S3	control	51 840	8 385	28 512	14 902	41	38	27	99.92	34 560	1.0
S3	sensing	51 840	7 176	27 328	17 293	43	36	27	99.92	34 560	1.0
S4	telemetry	345 600	24 331	292 656	28 553	60	52	32	99.98	–	–
S4	control	345 600	24 351	294 344	26 853	52	49	29	99.98	–	–
S4	sensing	345 600	23 497	286 394	35 657	52	45	28	99.98	–	–




Figure 4: Design-space funnel across scales and workloads (logarithmic scale).

6.3. RQ3: when pruning pays off

Using the measured overheads and evaluation counts in (9), the break-even evaluator cost is $1\text{--}3\mu\text{s}$ on S2 and below $1\mu\text{s}$ on S3, i.e. any evaluator more expensive than a few floating-point operations already justifies the pruning overhead. Fig. 6 shows the speed-up over exhaustive enumeration as a function of C_E . At the measured cost of the discrete-event evaluator (19.4ms) the predicted speed-up is $78\times$ on S2 and $558\times$ on S3, and it saturates at the ratio of evaluation counts, $81\times$ and $640\times$ respectively. With the cheap analytical evaluator the measured speed-up is only $1.3\text{--}3.3\times$, because then the bound computation dominates the runtime; this is exactly the regime in which pruning is *not* the right tool, and we report it rather than omit it.

6.4. RQ4: comparison with budget-matched exploration

Table 6 reports the comparison over 20 seeds. On S3 NSGA-II reaches a median Pareto recall of $0.50\text{--}0.74$ and an HV ratio of $0.94\text{--}0.98$ only at a budget of 1000 full evaluations, while random




Figure 5: Exact Pareto front recovered on S3/telemetry, projected on latency–cost and latency–energy; colour encodes the remaining objectives.

search never recovers a single exact front point at that scale. HCA-DSE attains recall, precision and HV ratio of 1.0 with 41 to 54 evaluations, i.e. between $18\times$ and $24\times$ fewer evaluations than the best budget of the stochastic baseline. At the smaller S2 scale NSGA-II is competitive in hypervolume (0.999 at $B = 1000$) but still misses 9–22% of the exact front points. Fig. 7 shows the seed-level distributions.

The comparison should be read as a statement about guarantees, not about the quality of evolutionary search: NSGA-II converges towards the front and would reach it with more budget, but it cannot certify that anything is missing, whereas HCA-DSE returns the exact front and does so at a small fraction of the evaluation budget — provided admissible bounds exist.

6.5. Ablation: what each layer contributes

Table 7 and Fig. 8 isolate the layers on S3. Structural pruning alone leaves the full 34 560 well-formed candidates to be evaluated. Feasibility bounds alone reduce this to 7 025–12 700 ($2.7\times$ – $4.9\times$). Dominance bounds alone are far more aggressive, leaving 327–415 evaluations, but the two are complementary: combined they leave 41 to 54, i.e. another $6\times$ – $10\times$ beyond dominance alone, because feasibility pruning removes exactly the sub-trees for which the dominance test would otherwise have to be repeated deep in the tree. All variants preserve the exact front.

6.6. RQ5: scalability

Table 8 and Fig. 9 extend the raw design space over five orders of magnitude, from 1.3×10^3 to 8.3×10^7 candidates. The number of visited partial states grows from 391 to 16 778, i.e. by a factor of 43 while the space grows by a factor of 64 000; the number of full evaluations is essentially flat at 30 to 61; the runtime stays below 2.4s and the peak memory below 90 KiB, because the Cartesian product is never materialised. The fraction of visited candidates falls from 0.30 to 1.7×10^{-4} .




Figure 6: Speed-up over exhaustive enumeration as a function of the evaluator cost C_E , from the measured overheads and evaluation counts.

6.7. RQ6: transfer to a real deployment

Against the discrete-event evaluator the analytical model attains Spearman’s ρ of 0.92, 0.92 and 0.96 for telemetry, control and sensing, with MAPE of 39%, 44% and 14%. On the containerised testbed (Table 9, Fig. 10) the pooled prediction error over the 36 architectures is 28.7% (MAPE), with a median absolute error of 7.5 ms, and the pooled rank correlation is $\rho = 0.97$, $\tau = 0.89$. The decision-relevant metric is the pairwise ordering agreement: among the 498 architecture pairs that the model separates by more than 25%, the measured order agrees in 98.2% of the cases; among the 561 pairs separated by more than 10%, in 96.8%.

Two systematic effects are visible. For the data-intensive sensing workload the model *over-*predicts latency by roughly 30% on eight of twelve architectures; because the bias is systematic rather than random, the rank correlation for that workload is the highest of the three ($\rho = 0.95$) while its MAPE is the worst (50%). And four architectures — all sensing designs with $\rho \geq 0.75$ — were not reproducible across repeats, varying by factors of two to five (for instance 157, 293 and 556 ms against a predicted 82 ms). Near saturation neither the analytical model nor the deployment itself is stable, which is a result about the limits of design-time prediction rather than measurement




Figure 7: Hypervolume ratio against the exhaustive front on S3 over 20 seeds. The red line marks the exact front recovered by HCA-DSE.

noise to be filtered out.

6.8. A model correction found by the measurements

The first version of the analytical model used an M/M/1 waiting term. Because stage execution is deterministic, that formulation overestimates waiting by a factor of two, and replacing it with the M/D/1 expression of (3) reduced the MAPE from 28.9% to 23.7% on an identical set of measurements. Interestingly the rank correlation was slightly *higher* for the M/M/1 variant (0.918 versus 0.889 on that testbed configuration), i.e. accuracy and ordering do not improve in lockstep — a practical reminder that a model tuned for absolute prediction is not automatically the better model for exploration.

7. Discussion

7.1. Why hierarchical pruning works here

The results of Section 6 are driven by a property of the design space rather than by a clever search order: in a heterogeneous edge-cloud CPS the vast majority of candidate architectures is either ill-formed, or violates a hard constraint that is already implied by a handful of early decisions, or is dominated by a design that is cheaper in every objective. On S3 only 34 560 of 51 840 candidates are even well-formed, of those a further 7 000 to 12 700 survive the feasibility bounds, and of those fewer than 60 need to be evaluated. Every candidate removed at those layers would otherwise consume a full simulation.

The ablation makes the mechanism explicit. Dominance pruning is the stronger of the two bound-based layers, but it is only effective once an archive of good designs exists; feasibility pruning removes precisely the sub-trees in which no such design will ever be found, so it reduces the number of tree nodes at which the dominance test has to be evaluated. The combination is roughly an order of magnitude stronger than either layer alone.




Figure 8: Full evaluations per pruning configuration on S3 (logarithmic scale).

7.2. When HCA-DSE is the right tool — and when it is not

The method is advantageous when the design space is large, the constraints are tight, the evaluator is expensive, and admissible bounds can be constructed. The crossover analysis makes the third condition quantitative: with the analytical evaluator, whose cost is comparable to that of the bounds themselves, the wall-clock gain is only $1.3\text{--}3.3\times$ despite a 99.9% reduction in evaluations, whereas with a discrete-event evaluator the same reduction turns into two orders of magnitude. Since industrial dCPS exploration reports evaluation costs of up to a minute per design point [1], the practically relevant regime is firmly on the favourable side of the crossover.

Conversely, the method offers little when the space is small, the constraints are loose, the evaluation is nearly free, or — most importantly — when no admissible bound can be constructed for the objectives of interest. If a bound cannot be proven admissible it must not be used: the guarantee of Theorem 1 rests entirely on that property, and an inadmissible bound silently converts an exact method into an approximate one without any indication in the output.

7.3. Relation to symbolic exact exploration

Symbolic approaches [4] prune partial designs with a constraint solver and are, in principle, more general than closed-form bounds: any constraint the solver can express can participate in the pruning. The price is that the objectives must be assignment monotonous for dominance checks on incomplete assignments to be sound. Our construction trades generality for applicability to contention-dependent objectives: Lemma 2 together with the decision ordering of Remark 1 recovers exactness for a latency model with queueing, at the cost of bounds that must be derived per objective. The two directions are complementary, and combining a solver-based front end with contention-aware bounds is a natural next step.




Figure 9: Visited partial states and full evaluations against the raw design-space size.

7.4. Relation to evaluator-centric methodologies

CompDSE [1] and HCA-DSE answer different questions. The former asks how a design point of an industrial dCPS can be modelled and evaluated faithfully; the latter asks which design points need not be evaluated at all. Their combination is straightforward in principle — the analytical model of Section 3 would serve as the bound generator and the trace-derived simulator as the evaluator — and it is the configuration in which our crossover analysis predicts the largest benefit.

7.5. Accuracy versus ordering

The empirical validation supports a narrower claim than “the model is accurate”. The model’s absolute error is 16–50% depending on the workload, which is in the same range as the 20% execution-time deviation reported for a design-time placement tool validated on a real deployment [12]. What the exploration actually requires, however, is that the model orders architectures correctly, and that claim is supported much more strongly: 98% agreement on the pairs the model separates by more than 25%. The complement is equally informative — within a single workload the rank correlation is lower, because after the M/D/1 correction many designs of one workload differ by less than the measurement noise. A model that cannot separate near-equivalent architectures is not a defect of the exploration method but an argument for the two-level scheme: analytical




Figure 10: Analytical latency against the measured median p_{50} latency of the same architecture on the containerised testbed. Open markers denote the four architectures whose repeats differ by a factor of 1.5 or more.

bounds remove the bulk of the space, and the surviving handful of candidates is where an expensive high-fidelity evaluator should be spent.

8. Threats to validity

Internal validity. The correctness of the method rests on the admissibility of the bounds. We address this at three levels: proofs (Lemmas 1–2, Propositions 1–3, Theorem 1); an exhaustive machine check of $LB_i(p) \leq f_i(x)$ over all prefixes and all completions on S1 for the three workloads; and a direct check that, for every pruned prefix, no completion contradicts the claim of the corresponding rule. We deliberately separate *theoretical* correctness from *implementation* correctness and verify both, since an admissible bound that is implemented incorrectly is indistinguishable, in the output, from an inadmissible one. Structural pruning is part of the definition of the design space rather than a consequence of the constraints, and the exhaustive baseline uses the same predicate; this was discovered by the pruning-safety test and is stated explicitly rather than absorbed into the comparison.

Construct validity. The analytical model abstracts a real deployment: M/D/1 waiting, no transport-protocol effects, a linear power model and normalised cost units. The testbed validation quantifies the resulting gap and shows that it affects absolute prediction much more than ordering. The sensing workload exposes the main structural gap — for large payloads over a shaped link the model does not capture transport-level effects and over-predicts systematically. The energy and cost models are not validated against measurements at all; they are used for comparative exploration only, and no absolute energy or cost claim is made.

External validity. The workloads are synthetic, although chosen to be qualitatively different rather than rescaled versions of one another, and the architecture template is fixed to four tiers with three stages. The empirical study is further restricted to configurations that fit on a single machine (≤ 4 edge nodes, ≤ 2 gateways), and link shaping is applied in the forward direction of each hop, matching the one-way transfers counted by the model. Results on a geographically distributed deployment may differ, particularly for the WAN hop.

Conclusion validity. Deterministic methods (exhaustive, HCA-DSE) require no statistical treatment. Stochastic baselines are run with 20 independent seeds and compared with non-parametric tests (Mann–Whitney U , Cliff’s δ). Testbed configurations are repeated three times and reported by their median; the four configurations whose repeats differ by more than a factor of 1.5 are reported separately rather than discarded, and all four operate at $\rho \geq 0.75$.

Limitations. The method assumes a static workload during one exploration, does not model run-time migration or autoscaling, and requires bounds to be derived per objective — a manual step that must be repeated when the objective set changes. No comparison against an exact symbolic or mixed-integer formulation is provided: the queueing term is non-linear in the decision variables, so a mixed-integer program would require linearisation and would therefore optimise a different model, but this leaves the question of relative efficiency against solver-based exact exploration open.

9. Conclusions

Architecture exploration for heterogeneous distributed cyber-physical systems is limited less by the size of the design space than by the cost of evaluating a single design point. This paper presented HCA-DSE, a hierarchical constraint-aware exploration method that removes structurally ill-formed, provably infeasible and provably dominated regions of the space before any full evaluation is performed. Its technical core is a family of admissible bounds for *contention-dependent* objectives, obtained by a term-wise relaxation combined with a decision ordering under which the queueing term of a tier becomes exactly computable on a partial design. This removes the assignment-monotonicity assumption on which existing exact exploration methods rely, and we proved that the exact Pareto front is preserved.

Across four design-space scales, three workload classes and spaces of up to 8.3×10^7 candidate architectures, HCA-DSE performed 26 to 61 full evaluations, a reduction of 99.90% to 99.98% with respect to the raw space, while reproducing the exhaustive Pareto front exactly wherever exhaustive enumeration is tractable. Both the admissibility of the bounds and the equality of the fronts were verified exhaustively by machine. Budget-matched NSGA-II required 1000 evaluations to reach a median Pareto recall of 0.50–0.91. The break-even evaluator cost is in the microsecond range, and with a discrete-event evaluator the measured speed-up reaches $558\times$. On a containerised

edge-gateway-cloud testbed with per-node link shaping, 36 architectures and 108 measured runs, the analytical ordering agreed with the measured ordering in 98% of the architecture pairs that the model separates by more than 25%.

The method is intended as design-time decision support: it does not require runtime adaptation, learned performance models, or a high-fidelity simulator in the inner loop. Its natural continuation is the combination with an evaluator-centric methodology, in which analytical bounds reduce an industrial design space to a few dozen candidates and a trace-derived simulator evaluates only those; and the derivation of tighter, possibly surrogate-assisted bounds that remain provably admissible.

Data availability

The complete implementation, all experiment scripts, the testbed and every result file required to regenerate the figures and tables of this paper are available in the replication package accompanying the submission.

References

- [1] M. Saadatmand, et al., CompDSE: A methodology for design space exploration of computing subsystems within complex cyber-physical systems, *IET Cyber-Physical Systems: Theory & Applications* (2025). doi:10.1049/cps2.70019.
- [2] M. Herget, M. Saadatmand, M. Bor, I. González Alonso, T. Stefanov, B. Akesson, A. D. Pimentel, Design space exploration for distributed cyber-physical systems: State-of-the-art, challenges, and directions, in: 25th Euromicro Conference on Digital System Design (DSD), 2022. doi:10.1109/DSD57027.2022.00068.
- [3] K. Deb, A. Pratap, S. Agarwal, T. Meyarivan, A fast and elitist multiobjective genetic algorithm: NSGA-II, *IEEE Transactions on Evolutionary Computation* 6 (2) (2002) 182–197. doi:10.1109/4235.996017.
- [4] K. Neubauer, P. Wanko, T. Schaub, C. Haubelt, Exact multi-objective design space exploration using ASPmT, in: *Design, Automation & Test in Europe Conference (DATE)*, 2018, pp. 257–260. doi:10.23919/DATE.2018.8342014.
- [5] C. Haubelt, K. Neubauer, T. Schaub, P. Wanko, Design space exploration with answer set programming, *KI - Künstliche Intelligenz* 32 (2018) 205–206. doi:10.1007/s13218-018-0530-3.
- [6] S. Petchrompo, D. W. Coit, A. Brintrup, A. Wannakrairot, A. K. Parlikad, A review of Pareto pruning methods for multi-objective optimization, *Computers & Industrial Engineering* 167 (2022) 108022. doi:10.1016/j.cie.2022.108022.
- [7] L. Mandow, J. L. Pérez-de-la Cruz, Multiobjective A* search with consistent heuristics, *Journal of the ACM* 57 (5) (2010) 1–25. doi:10.1145/1754399.1754400.
- [8] B. S. Stewart, C. C. White, Multiobjective A*, *Journal of the ACM* 38 (4) (1991) 775–814. doi:10.1145/115234.115368.
- [9] B. Kienhuis, E. F. Deprettere, P. van der Wolf, K. Vissers, A methodology to design programmable embedded systems: The Y-chart approach, *Embedded Processor Design Challenges* (2002) 18–37doi:10.1007/3-540-45874-3_2.
- [10] A. Varga, R. Hornig, An overview of the OMNeT++ simulation environment, *Simutools* (2008).
- [11] Y. Xiao, C. Oh, M. Lora, P. Nuzzo, ContrArc: Contract-based architecture exploration for cyber-physical systems, in: *Design, Automation & Test in Europe Conference (DATE)*, 2024.
- [12] H. Sedghani, F. Filippini, D. Ardagna, SPACE4AI-D: A design-time tool for AI applications resource selection in computing continua, *IEEE Transactions on Services Computing* (2024). doi:10.1109/TSC.2024.3479926.
- [13] J. Blank, K. Deb, pymoo: Multi-objective optimization in Python, *IEEE Access* 8 (2020) 89497–89509. doi:10.1109/ACCESS.2020.2990567.
- [14] E. Zitzler, L. Thiele, Multiobjective evolutionary algorithms: a comparative case study and the strength Pareto approach, *IEEE Transactions on Evolutionary Computation* 3 (4) (1999) 257–271. doi:10.1109/4235.797969.
- [15] H. Ishibuchi, H. Masuda, Y. Tanigaki, Y. Nojima, Modified distance calculation in generational distance and inverted generational distance, *Evolutionary Multi-Criterion Optimization (EMO)* (2015) 110–125doi:10.1007/978-3-319-15892-1_8.
- [16] L. Kleinrock, *Queueing systems, volume 1: Theory* (1975).

- [17] H. B. Mann, D. R. Whitney, On a test of whether one of two random variables is stochastically larger than the other, *The Annals of Mathematical Statistics* 18 (1) (1947) 50–60.
- [18] N. Cliff, Dominance statistics: Ordinal analyses to answer ordinal questions, *Psychological Bulletin* 114 (3) (1993) 494–509. doi:10.1037/0033-2909.114.3.494.
- [19] S. Hemminger, Network emulation with NetEm, *Linux Conf Au* (2005).

Table 6: Budget-matched comparison. Medians over 20 seeds for the stochastic methods.

scale	workload	method	evals	recall	HV ratio	IGD ⁺	time [s]
S2	telemetry	exhaustive	3 888	1.000	1.000	0.0000	0.059
		random ($B=100$)	100	0.037	0.520	0.2030	0.001
		random ($B=250$)	250	0.037	0.730	0.1244	0.002
		random ($B=500$)	500	0.111	0.853	0.0766	0.005
		random ($B=1000$)	1000	0.185	0.899	0.0540	0.011
		NSGA-II ($B=100$)	100	0.000	0.569	0.1889	0.013
		NSGA-II ($B=250$)	250	0.204	0.868	0.0519	0.039
		NSGA-II ($B=500$)	500	0.667	0.981	0.0101	0.091
		NSGA-II ($B=1000$)	1000	0.907	0.999	0.0005	0.207
		HCA-DSE	48	1.000	1.000	0.0000	0.044
S2	control	exhaustive	3 888	1.000	1.000	0.0000	0.056
		random ($B=100$)	100	0.000	0.522	0.2114	0.001
		random ($B=250$)	250	0.040	0.728	0.1248	0.002
		random ($B=500$)	500	0.120	0.846	0.0776	0.005
		random ($B=1000$)	1000	0.160	0.897	0.0530	0.012
		NSGA-II ($B=100$)	100	0.000	0.565	0.1573	0.014
		NSGA-II ($B=250$)	250	0.280	0.900	0.0398	0.041
		NSGA-II ($B=500$)	500	0.680	0.993	0.0065	0.095
		NSGA-II ($B=1000$)	1000	0.780	0.999	0.0030	0.234
		HCA-DSE	38	1.000	1.000	0.0000	0.044
S2	sensing	exhaustive	3 888	1.000	1.000	0.0000	0.041
		random ($B=100$)	100	0.000	0.201	0.3725	0.001
		random ($B=250$)	250	0.062	0.579	0.1720	0.002
		random ($B=500$)	500	0.125	0.804	0.0944	0.004
		random ($B=1000$)	1000	0.188	0.868	0.0732	0.009
		NSGA-II ($B=100$)	100	0.000	0.269	0.3144	0.014
		NSGA-II ($B=250$)	250	0.250	0.723	0.0983	0.046
		NSGA-II ($B=500$)	500	0.562	0.937	0.0224	0.110
		NSGA-II ($B=1000$)	1000	0.875	1.000	0.0017	0.270
		HCA-DSE	26	1.000	1.000	0.0000	0.024
S3	telemetry	exhaustive	34 560	1.000	1.000	0.0000	0.372
		random ($B=100$)	100	0.000	0.324	0.4165	0.001
		random ($B=250$)	250	0.000	0.403	0.2781	0.002
		random ($B=500$)	500	0.000	0.605	0.2188	0.004
		random ($B=1000$)	1000	0.000	0.653	0.1500	0.009
		NSGA-II ($B=100$)	100	0.000	0.069	0.5304	0.012
		NSGA-II ($B=250$)	250	0.000	0.294	0.2763	0.041
		NSGA-II ($B=500$)	500	0.258	0.754	0.1242	0.094
		NSGA-II ($B=1000$)	1000	0.500	0.945	0.0466	0.235
		HCA-DSE	54	1.000	1.000	0.0000	0.156
S3	control	exhaustive	34 560	1.000	1.000	0.0000	0.349
		random ($B=100$)	100	0.000	0.312	0.4466	0.001
		random ($B=250$)	250	0.000	0.377	0.3157	0.002
		random ($B=500$)	500	0.000	0.579	0.2286	0.004
		random ($B=1000$)	1000	0.000	0.652	0.1488	0.009
		NSGA-II ($B=100$)	100	0.000	0.114	0.4307	0.012
		NSGA-II ($B=250$)	250	0.037	0.539	0.2269	0.041
		NSGA-II ($B=500$)	500	0.315	0.877	0.0456	0.099
		NSGA-II ($B=1000$)	1000	0.648	0.983	0.0236	0.237
		HCA-DSE	41	1.000	1.000	0.0000	0.170
S3	sensing	exhaustive	34 560	1.000	1.000	0.0000	0.384
		random ($B=100$)	100	0.000	0.045	0.5926	0.001
		random ($B=250$)	250	0.000	0.224	0.4782	0.002
		random ($B=500$)	500	0.000	0.361	0.2993	0.004
		random ($B=1000$)	1000	0.000	0.476	0.2061	0.009
		NSGA-II ($B=100$)	100	0.000	0.023	0.7297	0.014
		NSGA-II ($B=250$)	250	0.000	0.126	0.3366	0.040
		NSGA-II ($B=500$)	500	0.185	0.682	0.2273	0.101
		NSGA-II ($B=1000$)	1000	0.741	0.944	0.1677	0.259
		HCA-DSE	23	1.000	1.000	0.0000	0.107

Table 7: Ablation on S3. For each workload: visited partial states, full evaluations and runtime.

variant	telemetry			control			sensing		
	states	evals	[ms]	states	evals	[ms]	states	evals	[ms]
structural only	56 797	34 560	424	56 797	34 560	400	56 797	34 560	355
structural + feasibility bounds	22 205	12 700	451	16 110	7 025	330	21 084	10 956	334
structural + dominance bounds	15 557	327	485	15 648	407	412	7 961	415	214
HCA-DSE (all layers)	5 832	54	166	5 826	41	147	4 393	43	118

Table 8: Scalability. For each workload: visited partial states, full evaluations and runtime [s].

$\mathcal{X}$	telemetry			control			sensing		
	states	evals	[s]	states	evals	[s]	states	evals	[s]
1.3×10^3	659	43	0.07	660	33	0.06	391	18	0.03
5.8×10^3	1 994	57	0.23	1 998	53	0.23	1 142	30	0.10
5.2×10^4	7 291	60	0.96	7 294	56	0.92	5 183	53	0.55
3.5×10^5	10 346	61	1.39	10 349	57	1.34	8 643	59	0.95
2.8×10^6	14 081	61	1.91	14 084	57	1.88	12 056	61	1.40
1.6×10^7	15 618	61	2.14	15 621	57	2.06	13 216	61	1.59
8.3×10^7	16 778	61	2.31	16 781	57	2.23	13 888	61	1.72

Table 9: Testbed validation. Pair agreement is the share of architecture pairs whose measured order matches the model order, among pairs the model separates by more than 10% resp. 25%.

subset	n	MAPE [%]	med. AE [ms]	Spearman ρ	Kendall τ	pair > 10%	pair > 25%	repeat spread [ms]
telemetry (all)	12	20.3	7.54	0.839	0.727	0.852	0.895	0.23
telemetry ($\rho \leq 0.75$)	6	16.2	2.91	0.829	0.733	0.818	1.0	0.18
control (all)	12	15.7	2.56	0.872	0.779	0.917	0.909	0.27
control ($\rho \leq 0.75$)	10	9.5	2.4	0.778	0.674	0.852	0.833	0.17
sensing (all)	12	50.0	20.0	0.949	0.87	0.957	0.957	2.17
sensing ($\rho \leq 0.75$)	10	49.7	19.28	0.985	0.944	1.0	1.0	1.86
sensing (reproducible)	8	41.7	18.77	0.994	0.982	1.0	nan	1.49
all workloads (all)	36	28.7	7.54	0.97	0.893	0.968	0.982	0.49
all workloads ($\rho \leq 0.75$)	26	26.5	3.91	0.948	0.861	0.962	0.982	0.4
all workloads (reproducible)	32	23.9	4.54	0.96	0.881	0.963	0.979	0.38